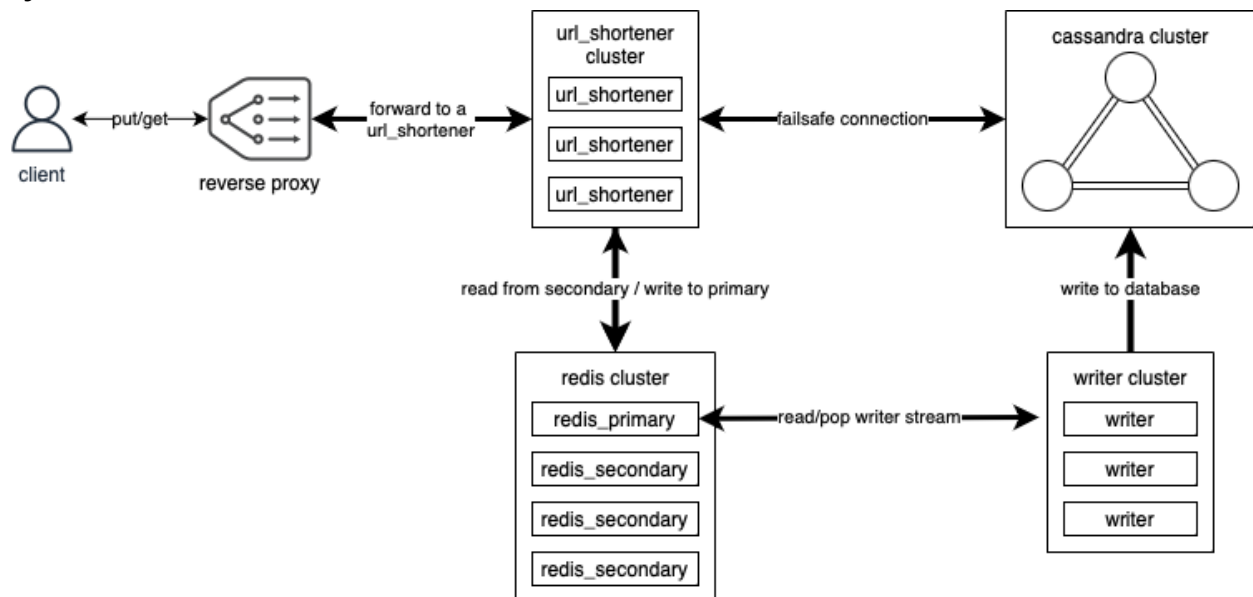


## System Architecture



reverse\_proxy:

- forwards the clients request to one of the url\_shorteners

url\_shortener:

- connects to redis\_primary, one redis\_secondary and the cassandra cluster
- parses the clients input
- reads from redis\_secondary
- if redis\_secondary is down, reads from redis\_primary
- if redis\_primary is down or does not contain required data, reads from cassandra
- writes to redis\_primary
- if redis\_primary is down, writes to cassandra

redis cluster:

- single primary and multiple secondaries
- primary contains a list for use by writers

writers:

- pop from redis\_primary list
- attempt writing to cassandra database until successful

cassandra\_cluster:

- persists data

dashboard:

- monitors all swarm services through docker

## System Discussion

### Cache

- Redis is used as a cache for short:long pairs (redis is also used for the writer stream)
- Uses LRU policy to evict keys
- Each short/long pair in the cache has a finite TTL to ensure consistency and fault tolerance

### Consistency

- Every request is timestamped to ensure that the latest PUT request is stored in the database.
- The redis cache (excluding the streaming list) has a TTL timer for each key value pair to ensure it becomes consistent with the cassandra database if redis malfunctions but is not restarted (therefore some requests are sent directly to cassandra but don't exist within redis)
- In the event that redis goes down, the redis cache (excluding streaming list) is purged as data in the cache may be outdated.

### Availability / Fault tolerance

- Multiple url\_shorteners so if one goes down, another is available to hand user requests.
- If redis\_secondaries go down, url\_shorteners will redirect all requests to redis\_primary
- If redis cluster goes down, url\_shorteners will redirect all requests to cassandra
- If cassandra is down, PUT requests will still succeed and be stored in a redis queue, writers will then send data from the queue to cassandra when it restarts. GET requests will succeed only if requested data is in the redis cache
- redis\_primary stream is persisted in case redis goes down
- Cassandra uses replication factor of 2 to ensure data recoverability if a node goes down

### Orchestration

- Single script starts the cassandra cluster and the system stack
- Everything other than cassandra is deployed from a single docker-compose file
- Adding more url\_shorteners, redis\_secondaries and writers is possible by changing the 'replicas' number in the docker-compose file
- Adding / Removing cassandra nodes can be done through the add/removeCassNode script

### Healing

- Every task in the swarm has a health check, if the health check fails or the process exits unexpectedly, docker will automatically restart the task.
- Cassandra uses replication factor of 2 to ensure data recoverability if a node goes down

### Load Balancing

- each url\_shortener attempts to connect to a different redis\_secondary through dockers load balancing

- traefik is used as reverse proxy to forward user requests to a different url\_shortener using round-robin

### Scalability

- Storage capacity can be horizontally scaled by adding additional cassandra nodes using script
- Throughput can be scaled by increasing the number url\_shortener replicas and in the docker-compose file
- Redis storage size can be vertically scaled by increasing max size in the docker-compose file
- Redis read throughput can be horizontally scaled by increasing redis\_secondary replicas in the docker-compose file

### Monitoring

- A dashboard UI accessed at URL: localhost:4001 shows the state of every service within the swarm including the dashboard itself
- Health checks ensure service states are up to date
- node status is used to monitor cassandra cluster status

### Logging

- All services output logs which are accessible through the Monitor dashboard
- Logs are persisted in each machine by docker

### Strengths

- high availability, some url\_shorteners, redis\_secondaries, redis\_primary and writers can go down and the system will remain functional
- easy to scale by changing the number of replicas in the docker-compose file

### Weaknesses


- To have additional reverse proxies, they must each publish different ports, therefore they must be configured manually.
- url\_shortener uses flask which is single threaded, therefore it does not utilize system resources efficiently, docker swarm does mitigate this by stacking services on the same system

## System Performance


### Setup

- 4 virtual machines each having 1 core and 2GB ram
- testing done on host machine using remaining 4 threads
- tests done using wrk '[GitHub - wg/wrk: Modern HTTP benchmarking tool](#)'
- architecture contains 1 writer, 2 url\_shorteners, 2 redis secondaries, 3 cassandra nodes, 1 redis primary.

**PUT request across 4 concurrent connections**




**GET request across 4 concurrent connections**



- both put and get requests scale linearly
- puts are significantly slower as all connections are directed to a single redis\_primary node
- average latency was below 4ms for all requests

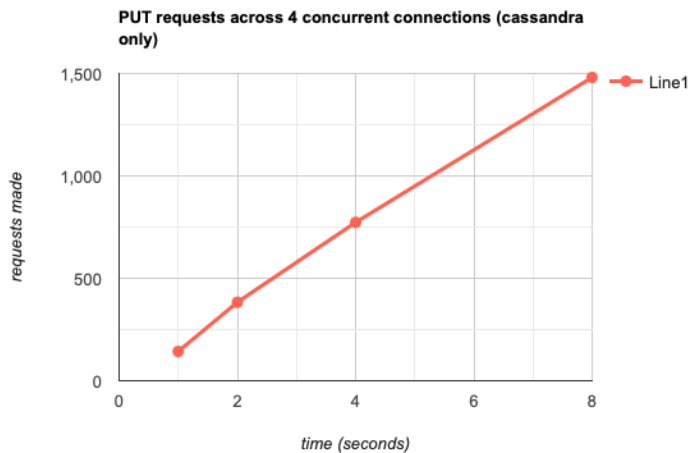
**GET requests across 4 concurrent connections (cassandra only)**



- get requests made directly to the cassandra were significantly slower, only completing 5000 in 8 seconds which is 6 times slower than get requests using redis with 2 secondaries

#### PUT Performance without writers


(same configuration as previous tests, PUT requests sent directly to cassandra)



Without the writer architecture, PUT requests suffer considerably, only managing to complete ~1500 requests in 8 seconds which is over 5 times slower than using redis with writers.

#### Notes:

- Increasing the number of writers does not improve PUT performance as the bottleneck is the single redis\_primary node. More writers can however improve the reliability and consistency of the service.
- Having more redis\_secondaries than url\_shorteners is inefficient as the url\_shorteners use flask which is single threaded therefore each url\_shortener can only utilize the performance of a single redis\_secondary at most. Therefore the number of redis\_secondaries should be  $\leq$  number of url\_shorteners
- only 2 url\_shorteners were used as the VMs only have a single core each therefore adding more url\_shorteners would result in more context switching resulting in worse performance. Using more VM nodes or having access to more cores would significantly improve performance as the main bottleneck for get requests (which exist in redis) is the url\_shorteners.



Using just 1 url\_shortener, performance dropped by between 30% and 60% as opposed to using 2 url\_shorteners.